

Inference of Cellular Developmental Time

transition-speed: default transition: fade autosize: false

“From Multivariate to Longitudinal Data” April 11, 2017

Caleb Lareau bit.ly/LareauBST245

Overview

- Motivation
- Single cell RNA-Seq
- Model Dataset
- EDA
- Methods of estimating developmental time
- PCA
- Probabilistic PCA; Bayesian PCA
- Gaussian Process Latent Variable Modeling

Trajectories...

=====

~3 million cells ~ 20 billion cells ~ 50 trillion cells

- What are the key points in development for disease?

=====

Questions from developmental biology

- What happens in a cell such that it becomes a brain, toe, or a heart? - When do these decisions get made? “Who” makes them? - How do the developmental trajectories of disease (leukemia / schizophrenia) differ from healthy individuals? - Can we identify important transition points and the genetic signature underlying them?

Waddington Landscape

Some cancers regain stemness programs

Stergachis *et al.*, Cell 2013

Stem cell-likeness in AML

Corces *et al.* Nature Genetics, 2016

How do we characterize single cells?

=====

Proserpio and Mahata, Immunology 2015

=====

=====

=====

=====

=====

=====

=====

=====

=====

EDA

```
class: small-code “{r, eval = FALSE} > dim(deng)
[1] 17585 255
      sum(deng == 0) / prod(dim(deng))
[1] 0.5019552
      head(sample(colnames(deng)))
[1] “earlyblast” “16cell” “4cell” “midblast” “lateblast” “16cell”
```

```

head(sample(rownames(deng)))
[1] "Gm7073" "Mir697" "Uqcrc2" "Ap2m1" "Slc10a3" "Ccr1"
““

```

=====

Linearly increasing

Dropoff

Linear Decreasing

Varying, no clear effect

V-shaped

Transition on/off

Sigmoidal with dropout

Overall picture

\forall gene g , fit OLS Regression with known timepoint t per cell–

$$\log_2(g + 1) \sim \beta_0 + \beta_1 t$$

Permuted

\forall gene g , fit regression with permuted timepoint t^* per cell

$$\log_2(g + 1) \sim \beta_0 + \beta_1 t^*$$

Permuted

\forall gene g , fit regression with permuted **factor** timepoint t^{**}

$$\log_2(g + 1) \sim \beta_0 + \beta_1 t^{**}$$

Statement of problem

Given a matrix \mathbf{Y} of D genes (features) by n samples, determine a latent vector \mathbf{P} with dimension $1 \times n$ that reflects the developmental trajectory of the n cells from the variance in D genes. - D can be thought of as a higher dimension space, and we want to infer d ($d < D$) latent variables in the gene data. - One of the d latent variables ideally reflects developmental ordering.

Perfect latent variable

class: small-code

```
“{r, eval = FALSE} > cor(runif(length(time)) %>% sort(), as.numeric(time)
%>% sort())^2
[1] 0.8799669 ““
```

PCA

Computing PCA

\mathbf{Y} is a $D \times n$ matrix. Compute the covariance matrix–

$$\mathbf{\Sigma} = E(\mathbf{Y}\mathbf{Y}^T) - \mu\mu^T$$

where $\mu = E(\mathbf{Y})$

Then compute the spectral decomposition of $\mathbf{\Sigma}$

$$\mathbf{\Sigma}a_j = \lambda_j a_j$$

for $j \in (1, \dots, D)$ Then a_j represent the eigenvectors of the data matrix \mathbf{Y} . Note the ordering of j is meaningful–

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_D \geq 0$$

Computing PCA

```
class: small-code “{r, eval = FALSE} > irlba::prcomp_irlba()
  prcomp()
  princomp()
““
```

PCA

```
class: small-code
“{r, eval = FALSE} > cor(pca$rotation[,1], as.numeric(time))^2
[1] 0.5933009 ““
```

Correlation with all PCs

Pause...

PCA by itself isn't satisfactory... Ideas for improvements?

Improving on PCA

1) Quantifying uncertainty 2) Non-linear latent variable inference

Uncertainty

From Buenrostro *et al.* bioRxiv 2017

Mannifold / Non-linear dimension learning

- Next several images taken from slides via Guy Wolf (Yale)

- These slides can be found here

=====

=====

=====

=====

Tackling 1 and building to 2

=====

=====

Slide from Neil Lawrence

Factor Analysis

\mathbf{Y} is a $D \times n$ matrix. \mathbf{x} is a $d \times n$ matrix ($d < D$). Want to relate \mathbf{x} and \mathbf{Y} and assume that the relationship is linear–

$$\mathbf{Y} = \mathbf{W}\mathbf{x} + \epsilon$$

where \mathbf{W} is a $D \times d$ matrix.

By convention (after locating the matrix),

$$\mathbf{x} \sim \mathcal{N}(0, \mathbf{I})$$

where \mathbf{I} is the identity matrix of dimension $d \times d$.

Factor Analysis II

Specify the error,

$$\epsilon \sim \mathcal{N}(0, \mathbf{\Psi})$$

where we assume $\mathbf{\Psi}$ to be a diagonal matrix $D \times D$. - Assuming the diagonal structure of $\mathbf{\Psi}$ means that all observational correlations is due to the latent variables. - In other words, Y_i for $i \in (1, \dots, n)$ are conditionally independent given \mathbf{x}

Factor Analysis III

Then Y_i for $i \in (1, \dots, n)$,

$$Y_i|x_i \sim \mathcal{N}(0, \Psi)$$

Then we can integrate out the latent variables–

$$p(Y_i|\mathbf{W}, \Psi) = \int p(Y_i|x_i, \mathbf{W}, \Psi)p(x_i)dx_i$$

Under this specification, we can write the MVN distribution for our observations–

$$\mathbf{Y} \sim \mathcal{N}(0, \mathbf{C}), \mathbf{C} = \mathbf{W}\mathbf{W}^T + \Psi$$

Computing PCA via Factor Analysis

Young-Whittle Factor Analysis (1950s) ψ_i element of the diagonal of Ψ ; add constraint $\psi_i = \sigma^2$ Assume σ^2 known, MLE yields same W as PCA

$$\mathcal{L} = \frac{-N}{2} \{D \log(2\pi) + \log |\mathbf{C}| + \text{tr}(\mathbf{C}^{-1}\mathbf{\Sigma})\}$$

$$\mathbf{C} = \mathbf{W}\mathbf{W}^T + \Psi$$

$$\mathbf{x} = \mathbf{C}^{-1}\mathbf{W}^T\mathbf{Y}$$

Probabilistic PCA (Tipping and Bishop)

- Assuming σ^2 may not be reasonable; want to estimate it from the data (keep in likelihood)
- Can estimate \mathbf{W}, σ^2 using EM & with a prior over \mathbf{x}

For \mathbf{W}_{MLE} ,

$$\sigma_{MLE}^2 = \frac{1}{D-d} \sum_{j=d+1}^D \lambda_j$$

- Similarly, we can integrate over \mathbf{W} given a prior, yielding

$$\mathbf{x} \sim \mathcal{N}(\mathbf{C}^{-1}\mathbf{W}^T\mathbf{Y}, \sigma^2\mathbf{C}^{-1})$$

Bayesian PCA

Bayesian PCA

- α_i represents the inverse of the variance
- for large α_i , contribution of latent factor i is low.
- Solved by EM / similar algorithm

Bayesian / Probabilistic PCA in R

class: small-code “{r, eval = FALSE} library(pcaMethods)

pca()

resPPCA <- pca(data, method="ppca", center=FALSE, nPcs=5) resBPCA <-
pca(data, method="bpca", center=FALSE, nPcs=5)

““

=====

Non-linear

$$\mathbf{Y} \sim \mathcal{N}(0, \mathbf{C}), \mathbf{C} = \mathbf{W}\mathbf{W}^T + \Psi$$

$$\mathbf{C}(\mathbf{W}_i, \mathbf{W}_j) = \mathbf{W}_i^T \mathbf{W}_j + \sigma^2 \delta_{ij}$$

$$\mathbf{C}(\mathbf{W}_i, \mathbf{W}_j) = \theta_{rbf} \exp\left(\frac{-\gamma}{2}(\mathbf{W}_i - \mathbf{W}_j)^T(\mathbf{W}_i - \mathbf{W}_j)\right) + \dots$$

- (2) being a special case (linear, iid) of (3)
- Computationally more challenging, but there are fast algorithms out there

GPLVM

GPLVM

class: small-code “{r, eval = FALSE} library(pseudogp)

fit <- fitPseudotime(data, smoothing_alpha = 30, smoothing_beta = 6, iter =
1000, chains = 1)

posteriorBoxplot(fit) ““

GPLVM – Noteable application

GPLVM

```
class: small-code
“{r, eval = FALSE} > cor(gplvm_means, as.numeric(time))^2
[1] 0.783522 “
```

GPLVM versus PCA 1

```
class: small-code
“{r, eval = FALSE}
  cor(pca$rotation[,1], as.numeric(time))^2
[1] 0.5933009
  cor(gplvm_means, as.numeric(time))^2
[1] 0.783522 “
```

GPLVM versus PCA 2

Wrapping up...

- GPLVM provide both a probabilistic and non-linear latent variable inference structure - All other published algorithms provide point estimates, difficult to ascertain uncertainty - Under this framework, Bayesian priors are allowed, which have been useful in published studies - Tools and methods useful in many data analyses contexts (including machine learning) - Motivated here by single cell analysis

Detailed look at early embryo neurogenesis?

=====

Thanks!